CPSC 340 Notation Guide

January 30, 2024

A guide to notation used in the course. Let me know if things are missing from this document that are not obvious.

Part 1: Supervised Learning

Throughout the course, we use n as the number of training examples and d as the number of the features. We use i when indexing a training example and j when indexing a feature.[1] We use x_{ij} as feature j in training example i, and we use y_i as the label of example i (so there are n values y_1, y_2, ..., y_n). We use y as a list of the n class labels, containing the label y_i in position i. We use x_i as the list of all features for example i, so x_i has d elements x_{i1}, x_{i2}, ..., x_{id} (and there are n lists x_1, x_2, ..., x_n). We use X as an $n \times d$ matrix containing all the features, so x_{ij} is element (i, j) of X and x_i gives the elements of row i of X. We use x^j to refer to all elements of column j, which is the list of values of feature j across all the n training examples.

Throughout the course, we use t as the number of test examples, and $\tilde{X}$ refers to a $t \times d$ matrix containing the test features. The notation $\tilde{x}_i$ refers to the features of test example i, whle $\tilde{x}_{ij}$ refers to feature j in test example i. We use $\tilde{y}$ as the true labels of the test examples, and $\tilde{y}_i$ as the label of test example i. We use $\hat{y}_i$ as the prediction of a model on example i, whether the prediction is made on training data or validation or test data (it should be obvious or not relevant from context).

When discussing validation sets, X_{train} and y_{train} are used as the subsets that we train on, while X_{validate} and y_{validate} are used as the subsets that we validate on. We use E to denote a generic prediction error, and usually this is followed with a subscript. For example, E_{train} is the training error, E_{test} is the test error, E_{gap} is the generalization gap.

We use c as a class label, and occasionally use n_c as the number of training examples in class c. We use ℓ as the number of possible class labels. We use the letter k generically throughout the course as something we count (typically a hyper-parameter of a model), and ϵ as a generic number that we want to be small.

Some method-specific notations used in this section:

- We use t as a particular decision stump threshold, and k as the number of thresholds.

- $p(y_i = \text{``}spam''|x_i)$ is used for the probability that the label y_i takes the value "spam" given that the features are x_i.

- $p(y_i|x_i)$ is used for the probability that the label is y_i given that the features are x_i (for example, y_i could be "spam" or "not spam" but without specifying a particular value).

- In the naive Bayes section, we're a little sloppy in that we use the same notation for the MLE on the training data and the true population value.

---

[1] When talking about two training examples, we sometimes use j as the index of the second training example.

- Naive Bayes uses n_{cjk} as the number of times that feature j is equal to k and the class label is c.

- Naive Bayes uses $p(\text{hello}|\text{spam})$ as short for $p(x_{ij} = \text{"hello"} \mid y_i = \text{"spam"})$.

- Decision theory slides use $\text{cost}(\hat{y}_i, \tilde{y}_i)$ as the cost of prediction $\hat{y}_i$ when the true label is $\tilde{y}_i$.

- In the norm slides we use r as a generic vector.

- We use $\|r\|_2$ as the L2-norm (square root of sum of squared the elements in the vector), $\|r\|_1$ as the L1-norm (sum of absolute values), and $\|r\|_\infty$ as the L∞-norm (max of absolute values). If the number is omitted, as in $\|w\|$, it refers to the L2-norm.

- For ensemble methods we use m as the number of models.

Part 2: Unsupervised Learning

We use $\hat{y}_i$ as the cluster predicted for example i and $\hat{y}$ as the set of predicted clusters for all n training examples. We use C as the set of indices of examples assigned to cluster c.

We use W as a k by d matrix where row c contains mean c. We use w_c as mean c, $w_{\hat{y}_i}$ to refer to the mean of the cluster of example i, and w_{cj} to refer to feature j in mean c. We use w^j as column j of the matrix W. We use $\hat{X}$ as predicted values of the matrix X, and similarly $\hat{x}_i$ are predicted values of x_i and $\hat{x}_{ij}$ are predicted values of x_{ij}.

We use μ as the mean of the data (with μ_j being the mean for feature j if we have more than one feature) and σ as the standard deviation (with σ_j being the standard deviation for feature j if we have more than one feature).

Part 3: Linear models

In this section we start treating x_i and y_i as vectors (instead of lists), so we now have to be careful about whether vectors are row-vectors or column-vectors. Our default choice is that everything is a column-vector, so each x_i is a $d \times 1$ vector and y is an $n \times 1$ vector. Since x_i is now a column-vector, we need to be careful to define row i of X as x_i^T (instead of just x_i).

We use w as the $d \times 1$ vector of regression weights. We normally index into w using w_j. We sometimes add a y-intercept ("bias") variable and use w_0 to denote this variable (in some settings later in the course β is used instead of w_0).

We use $\nabla f(w)$ to denote the gradient of a function f with respect to w. Assuming w has length d, this is a $d \times 1$ vector where position j contains the partial derivative of f with respect to w_j. We use r as the vector of "residuals", $r = Xw - y$. An individual element i of r would be $r_i = w^T x_i - y_i$.

Gradient descent uses w^t as the parameter vector on iteration t (so t has a separate meaning than "number of test examples" here). The distinction between w^t (iteration t of gradient descent) and w^j (column j of matrix W) should be clear from the context. We use α^t as the step size on iteration t. We use w^* as a minimum of $f(w)$. Stochastic gradient uses f_i to refer to the loss function on example i.

We use Z as an $n \times k$ matrix of features obtained under a change of basis, and z_i as the list of k features in the new basis for example i. When we do linear regression under a change of basis, we use v as the $k \times 1$ vector of parameters (instead of the usual $d \times 1$ vector w). We use $\tilde{Z}$ as the transformation of test data $\tilde{X}$.

We use λ as the (scalar) regularization parameter. It is assumed to be non-negative (and will almost always be positive).

We use $\text{sign}(\alpha)$ as a function that return $+1$ if α is positive and -1 if α is negative. We use o_i as the continuous prediction made on example i by binary classification models like logistic regression and SVM, so the final classification is given by $\text{sign}(o_i)$.

Multi-class classification uses the same matrix W as we used for k-means, and we use w_{y_i} as the w_c value for the true label y_i. We use o_{ic} for the continuous prediction of class c for example i, so the final classification on example i is given by $\arg\max c\{o_{ic}\}$.

We use $h(r)$ as the sigmoid function applied element-wise to a vector r.

Some method-specific notation used in this section:

- p is used as the degree of the polynomial in the polynomial basis, and we sometimes use Z_p when we want to specify specifically that we've used a degree-p basis.

- When we introduce convolutions we use x as signal, w as a filter, and z as the output of the filter.

- We use K as the $n \times n$ Gram matrix, containin $z_i^T z_j$ in position (i, j). We use $\tilde{K}$ as the $t \times n$ matrix containing $\tilde{z}_i^T z_j$ in position (i, j). We use u as the $n \times 1$ parameter vector when doing kernel methods for lienar models. The kernel function is written as $k(x_i, x_j)$.

- When introducing MLE/MAP, we use D as generic data (indexed by D_i if it splits into IID training examples), w as generic parameters, and $\hat{w}$ as the predicted MLE or MAP value of w.

Part 4: Latent-Factor Models

Linear latent-factor models use the approximation $X \approx ZW$, where we use the same notation for Z and W as above: Z is $n \times k$ with z_i^T as the rows and z_{ic} as individual elements, W is $k \times d$ with w_c^T as the rows and w^j as the columns and w_{cj} as the individual elements. To avoid expressions like $(w^j)^T z_i$, for inner products in this section we sometimes use notation $\langle w, x \rangle$ to represent the inner product $w^T x$. We use $\|X\|_F$ as the Frobenius norm (square root of sum of elements squared).

Part 5: Neural Networks

This section continues using the same notation, but we now use $W^{(2)}$ and $Z^{(2)}$ as the values in layer 2. We also use w_{c0} as the bias on hidden unit c, p as the index of a specific layer, m as the number of layers, and T as the length of a sequence (in RNNs/transformers).